

ECE30030/ITP30010 Database Systems

SQL DDL

Reading: Chapter 3

Charmgil Hong

charmgil@handong.edu

Spring, 2025

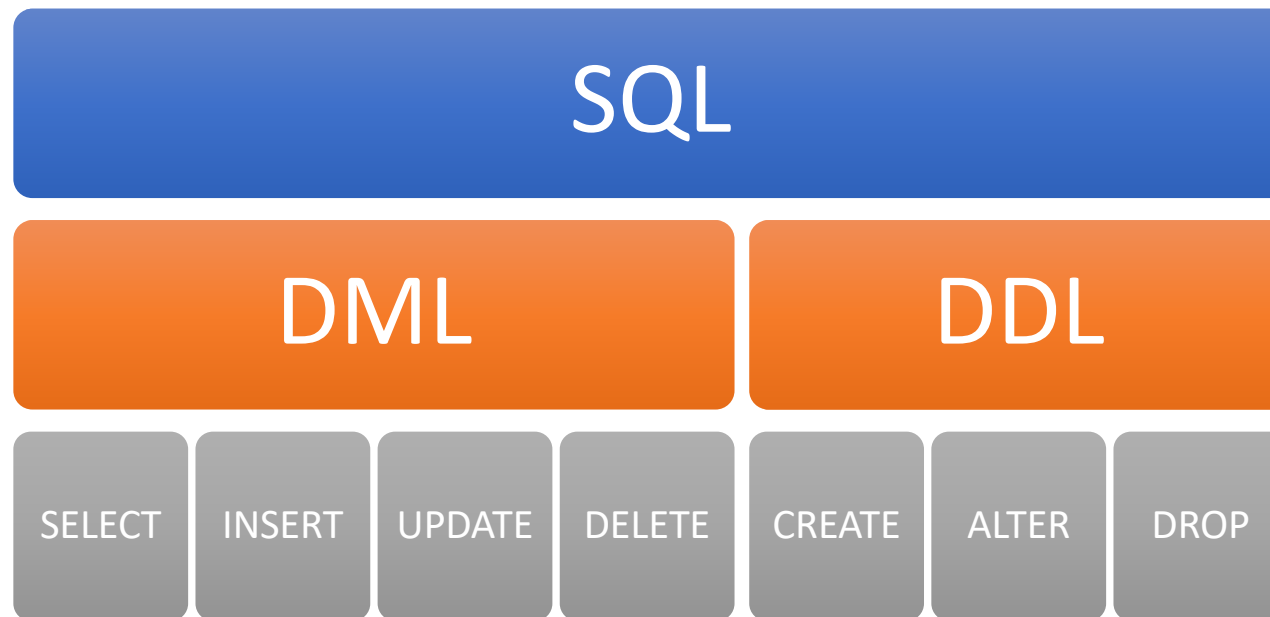
Handong Global University



Agenda

- SQL DDL (Data Definition Language)

SQL Commands



Data Definition Language

- The SQL data-definition language (DDL) allows the specification of information about relations, including:
 - The **schema** for each relation
 - The **type** of values associated with each attribute
 - The Integrity **constraints**
 - The set of **indices** to be maintained for each relation
 - Security and authorization information for each relation
 - The physical storage structure of each relation on disk

Domain Types in SQL

- SQL Data Types
 - **CHAR(*n*)**: Fixed length character string, with user-specified length *n*
 - Maximum length *n* = [0, 255]
 - **VARCHAR(*n*)**: Variable length character strings, with user-specified maximum length *n*
 - Maximum length *n* = [0, 65,535]
 - *If the length is always the same, use a CHAR-type attribute;
if you are storing wildly variable length strings, use a VARCHAR-type attribute*
 - **TEXT**: for strings longer than the range of VARCHAR
 - TINYTEXT 0 – 255 bytes
 - TEXT 0 – 65,535 bytes
 - MEDIUMTEXT 0 – 16,777,215 bytes
 - LONGTEXT 0 – 4,294,967,295 bytes

Domain Types in SQL

- Difference between CHAR and VARCHAR

Value	CHAR(4)	Storage	VARCHAR(4)	Storage
''	' '	4 bytes	''	1 bytes
'ab'	'ab '	4 bytes	'ab'	3 bytes
'abcd'	'abcd'	4 bytes	'abcd'	5 bytes
'abcdefg'	'abcd'	4 bytes	'abcd'	5 bytes

Domain Types in SQL

- “\”%ab%\””

Domain Types in SQL

- SQL Data Types
 - **INT, INTEGER:** *Integer* (a finite subset of the integers that is machine-dependent)
 - **SMALLINT:** *Small integer* (a machine-dependent subset of the integer domain type)
 - **BIGINT:** *Small integer* (a machine-dependent subset of the integer domain type)
- **TINYINT** and **MEDIUMINT** are also available

Domain Types in SQL

- Different R-DBMSs support different combinations of those integer types

	Bytes	MySQL	MS SQL	PostgreSQL	DB2
TINYINT	1	✓	✓		
SMALLINT	2	✓	✓	✓	✓
MEDIUMINT	3	✓			
INT/INTEGER	4	✓	✓	✓	✓
BIGINT	8	✓	✓	✓	✓

- *C.f.*, Oracle only has a NUMBER datatype

Domain Types in SQL

- SQL Data Types
 - **NUMERIC(p,d)**: **Fixed point number** (exact value) with user-specified precision of p digits, with d digits to the right of decimal point
 - *E.g.*, **NUMERIC**(3,1) allows 44.5 to be stored exactly, but not 444.5 or 0.32)
 - In MySQL, **DECIMAL** is NUMERIC
 - **FLOAT**: Floating point number (approximate) with **single-precision**
 - **REAL, DOUBLE**: Floating point number (approximate) with **double-precision**

Domain Types in SQL

- DECIMAL vs INT/FLOAT/DOUBLE
 - FLOAT and DOUBLE are faster than DECIMAL
 - DECIMAL values are exact
 - Example

floats: FLOAT	decimals: DECIMAL(3,2)
1.1	1.10
1.1	1.10
1.1	1.10

- SELECT SUM(...) → DECIMAL values are precise

SUM(floats)	SUM(decimals)
3.3000000715255737	3.30

Domain Types in SQL

- SQL Data Types
 - **DATE**: 'YYYY-MM-DD'
 - Range: 1000-01-01 to 9999-12-31
 - *E.g.*, '2020-03-01' for March 1, 2020
 - **TIME**: 'HH:MM:SS'
 - Range: -838:59:59 to 838:59:59
 - *E.g.*, '14:30:03.5' for 3.5 seconds after 2:30pm
 - **DATETIME**: 'YYYY-MM-DD HH:MM:SS'
 - Range: 1000-01-01 00:00:00 to 9999-12-31 23:59:59
 - **YEAR**: 'YYYY'
 - Range: 1901 to 2155, or 0000 (illegal year values are converted to 0000)

Domain Types in SQL

- SQL Data Types
 - **TIMESTAMP(*n*)**: **Unix time** (time since Jan 1, 1970)
 - Range: 1970-01-01 00:00:01 UTC to 2038-01-19 03:14:07 UTC
 - Typically used for logging (keeping records of all the system events)
 - Depending on size *n*, the display pattern changes

	Format
TIMESTAMP(14)	YYYYMMDDHHMMSS
TIMESTAMP(12)	YYMMDDHHMMSS
TIMESTAMP(10)	YYMMDDHHMM
TIMESTAMP(8)	YYYYMMDD
TIMESTAMP(6)	YYMMDD
TIMESTAMP(4)	YYMM
TIMESTAMP(2)	YY

Domain Types in SQL

- SQL Data Types
 - **BINARY(*n*)**: binary byte data type, with user-specified length *n*
 - Contains a byte strings (rather than a character string)
 - Maximum length $n = [0, 255]$
 - **VARBINARY(*n*)**: binary byte data type, with user-specified maximum length *n*
 - Maximum length $n = [0, 65,535]$
 - **BLOB**: Binary Large OBject data type
 - TINYBLOB 0 – 255 bytes
 - BLOB 0 – 65,535 bytes (65 KB)
 - MEDIUMBLOB 0 – 16,777,215 bytes (16 MB)
 - LONGBLOB 0 – 4,294,967,295 bytes (4 GB)

CREATE TABLE Construct

- A new relation is defined using the **CREATE TABLE** command:

```
CREATE TABLE r  
    (A1 D1, A2 D2, ..., An Dn,  
    (integrity-constraint1),  
    ...,  
    (integrity-constraintk))
```

- *r* is the name of the relation
 - Each *A_i* is an attribute name in the schema of relation *r*
 - Each *D_i* is the data type of values in the domain of attribute *A_i*
-
- Example: **CREATE TABLE** instructor(
 ID CHAR(5),
 name VARCHAR(20),
 dept_name VARCHAR(20),
 salary NUMERIC(8,2))

Integrity Constraints in CREATE TABLE

- SQL prevents any update to the database that violates an **integrity constraint**
 - Integrity constraints allow us to specify what data makes sense for us
- Types of integrity constraints
 - Primary key: **PRIMARY KEY** (A_1, \dots, A_n)
 - Foreign key: **FOREIGN KEY** (A_m, \dots, A_n) **REFERENCES** r
 - Unique key: **UNIQUE**
 - Not null: **NOT NULL**

- Example:

```
CREATE TABLE instructor(  
    ID                CHAR(5),  
    name              VARCHAR(20) NOT NULL,  
    dept_name         VARCHAR(20)  
    salary           NUMERIC(8, 2),  
    PRIMARY KEY (ID),  
    FOREIGN KEY (dept_name) REFERENCES department);
```


Declaring Keys

- An attribute or list of attributes may be declared as PRIMARY KEY or UNIQUE
 - Meaning: no two tuples of the relation may agree in all the attribute(s) on the list
 - That is, the attribute(s) do(es) **not allow duplicates** in values
 - PRIMARY KEY/UNIQUE can be used as an **identifier for each row**
 - Comparison: PRIMARY KEY vs UNIQUE

PRIMARY KEY	UNIQUE
Used to serve as a unique identifier for each row in a relation	Uniquely determines a row which is not primary key
Cannot accept NULL	Can accept NULL values (some DBMSs accept only one NULL value)
A relation can have only one primary key	A relation can have more than one unique attributes
Clustered index	Non-clustered index

Declaring Keys

- **CREATE TABLE** *student* (
 ID **VARCHAR(5),**
 name **VARCHAR(20) NOT NULL,**
 dept_name **VARCHAR(20),**
 tot_cred **NUMERIC(3,0),**
 PRIMARY KEY (*ID*),
 FOREIGN KEY (*dept_name*) **REFERENCES** *department*);

More Examples

- **CREATE TABLE** *student* (
 ID **VARCHAR(5) PRIMARY KEY,**
 name **VARCHAR(20) NOT NULL,**
 dept_name **VARCHAR(20),**
 tot_cred **NUMERIC(3,0),**
 FOREIGN KEY (*dept_name*) **REFERENCES** *department*);

More Examples

- **CREATE TABLE** *takes* (
 ID **VARCHAR**(5),
 course_id **VARCHAR**(8),
 sec_id **VARCHAR**(8),
 semester **VARCHAR**(6),
 year **NUMERIC**(4,0),
 grade **VARCHAR**(2),
 PRIMARY KEY (*ID, course_id, sec_id, semester, year*),
 FOREIGN KEY (*ID*) **REFERENCES** *student*,
 FOREIGN KEY (*course_id, sec_id, semester, year*)
 REFERENCES *section*);

More Examples

- **CREATE TABLE** *course* (
 course_id **VARCHAR**(8),
 title **VARCHAR**(50),
 dept_name **VARCHAR**(20),
 credits **NUMERIC**(2,0),
 PRIMARY KEY (*course_id*),
 FOREIGN KEY (*dept_name*) **REFERENCES** *department*);

More Examples

- **CREATE TABLE** *course* (
 course_id **VARCHAR**(8),
 title **VARCHAR**(50),
 dept_name **VARCHAR**(20) **DEFAULT** 'Comp. Sci',
 credits **NUMERIC**(2,0),
 PRIMARY KEY (*course_id*),
 FOREIGN KEY (*dept_name*) **REFERENCES** *department*);

More Examples

- **CREATE TABLE** *neighbors*(
 name **CHAR(30) PRIMARY KEY**,
 addr **CHAR(50) DEFAULT '123 Sesame St.'**,
 phone **CHAR(16));**
- Inserting Elmo is a neighbor:
 - **INSERT INTO** *neighbors* (*name*)
 VALUES ('Elmo');

name	addr	phone
'Elmo'	'123 Sesame St.'	NULL

More Examples

- **CREATE TABLE** *neighbors*(
 name **CHAR(30) PRIMARY KEY**,
 addr **CHAR(50) DEFAULT '123 Sesame St.'**,
 phone **CHAR(16) NOT NULL**);
- Inserting Elmo is a neighbor:
 - **INSERT INTO** *neighbors* (*name*)
 VALUES ('Elmo');
 - ➔ If *phone* were NOT NULL, this insertion would have been **rejected**

Table Updates (Updating Tuples)

- INSERT
 - **INSERT INTO** *instructor* **VALUES** ('10211', 'Smith', 'Biology', 66000)
- DELETE
 - **DELETE FROM** *student*
 - Remove all tuples from the *student* relation

Table Updates (Updating Table Schemas)

- DROP TABLE
 - **DROP TABLE r**
 - Remove relation r
- ALTER
 - **ALTER TABLE r ADD A D**
 - A is the name of the **new attribute** to add to relation r ; D is the **domain** of A
 - All **existing tuples in the relation are assigned *null*** as the value for the new attribute
 - **ALTER TABLE r DROP A**
 - A is the name of an **attribute** in r
 - Dropping of attributes not supported by many databases (MySQL does)

Table Updates (Updating Table Schemas)

- Examples
 - **DROP TABLE** time_slot_backup;
 - **ALTER TABLE** time_slot_backup **ADD** remark VARCHAR(20);
 - **ALTER TABLE** time_slot_backup **DROP** remark;

EOF

- Coming next:
 - Designing a database