

ECE30030/ITP30010 Database Systems

E-R Model

Reading: Chapter 6

Charmgil Hong

charmgil@handong.edu

Spring, 2025

Handong Global University



Agenda

- Designing a database
- E-R diagrams

Design Phases

- Initial phase: characterize fully the **data needs** of the prospective database users
- Second phase: choose a **data model**
 - Apply the concepts of the chosen data model
 - Translate the requirements into a **conceptual schema** of the database
 - A fully developed conceptual schema indicates the **functional requirements** of the enterprise
 - Describe the kinds of **operations** (or transactions) that will be performed on the data

Design Phases

- Final Phase: Move from an abstract data model to the **implementation** of the database
 - Logical Design – Deciding on the **database schema**
 - Database design requires that we find a “good” collection of relation schemas
 - *Business decision – What attributes should we record in the database?*
 - *Computer Science decision – What relation schemas should we have and how should the attributes be distributed among the various relation schemas?*
 - Physical Design – Deciding on the **physical layout** of the database

Design Phases

- In designing a database schema, we must ensure that we avoid two major pitfalls:
 - **Redundancy**: a bad design may result in repeated information
 - Redundant representation of information may lead to **data inconsistency among the various copies** of information
 - **Incompleteness**: a bad design may make certain aspects of the enterprise **difficult or impossible to model**
- Avoiding bad designs is not enough. There may be a large number of good designs from which we must choose

Design Approaches

- Entity Relationship Model
 - Models an enterprise as a collection of *entities* and *relationships*
 - Entity: a “thing” or “object” in the enterprise that is distinguishable from other objects
 - Described by *a set of attributes*
 - Relationship: an association among several entities
 - Represented diagrammatically by an *entity-relationship diagram* (E-R diagram)
- Normalization Theory
 - Formalize what designs are bad, and test for them

Agenda

- Designing a database
- **E-R diagrams**
 - Mapping cardinalities
 - Primary keys in E-R models
 - Weak entity sets
 - Reduction to relation schemas

E-R Model for Database Modeling

- The E-R data model was developed to facilitate database design by allowing specification of a **database schema**
 - Database schema represents the overall logical structure of a database
- The E-R data model employs three basic concepts:
 - Entity sets
 - Relationship sets
 - Attributes
- The E-R model has an associated **diagrammatic representation**
 - **E-R diagram** can express the **overall logical structure of a database graphically**

Entity Sets

- An **entity** is an object that exists and is distinguishable from other objects
 - *E.g.*, specific person, company, event, plant
- An **entity set** is a set of entities of the same type that share the same properties
 - *E.g.*, set of all persons, companies, trees, holidays
- An entity is represented by **a set of attributes**; *i.e.*, descriptive properties possessed by all members of an entity set
 - *E.g.*, *instructor* = (*ID*, *name*, *salary*)
course = (*course_id*, *title*, *credits*)
- A subset of the attributes form a **primary key** of the entity set; *i.e.*, **uniquely identifying each member of the set**

Representing Entity Sets in E-R Diagrams

- Entity sets can be represented graphically as follows:
 - Rectangles represent entity sets
 - Attributes listed inside entity rectangle
 - Underline indicates primary key attributes



Relationship Sets

- A **relationship** is an association among several entities

- *E.g.*,

44553 (Peltier)	<u>advisor</u>	22222 (<u>Einstein</u>)
<i>student</i> entity	relationship set	<i>instructor</i> entity

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where (e_1, e_2, \dots, e_n) is a relationship

- *E.g.*, $(44553, 22222) \in \text{advisor}$

Example: Entity and Relationship Sets

- Entity Sets – *instructor* and *student*

76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

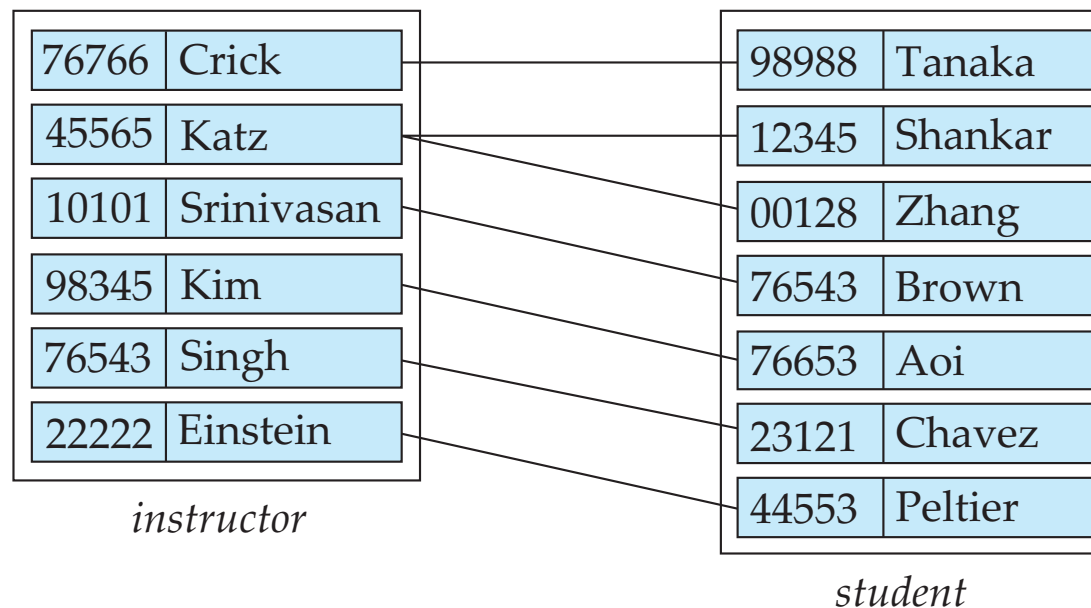
instructor

98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

student


Example: Entity and Relationship Sets

- Relationship Sets – define the relationship set *advisor* to denote the associations between students and the instructors who act as their advisors




Representing Relationship Sets via E-R Diagrams

- Diamonds represent relationship sets




Example: Entity and Relationship Sets

- An attribute can also be associated with a relationship set
 - *E.g.*, the *advisor* relationship set between entity sets *instructor* and *student* may have the attribute *date* which tracks when the student started being associated with the advisor




Relationship Sets with Attributes

- An attribute can also be associated with a relationship set



Roles

- Entity sets of a relationship need not be distinct
 - Each occurrence of an entity set plays a “role” in the relationship
 - *E.g.*, The labels “*course_id*” and “*prereq_id*” are called **roles**




Degree of a Relationship Set

- Binary relationship
 - Involves two entity sets (or degree two)
 - Most relationship sets in a database system are binary
- Relationships between more than two entity sets are rare but possible
 - *E.g., students* work on research *projects* under the guidance of an *instructor*
 - Relationship *proj_guide* is a ternary relationship between *instructor*, *student*, and *project*

Non-binary Relationship Sets

- Most relationship sets are binary
- There are occasions when it is more convenient to represent relationships as non-binary
- E-R diagram with a **ternary relationship**:




Complex Attributes

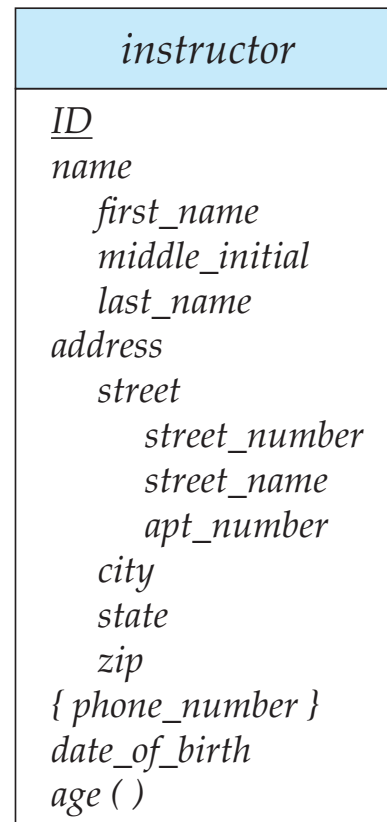
- Attribute types:
 - Simple and composite attributes
 - Single-valued and multivalued attributes
 - *E.g.*, multivalued attribute: *phone_numbers* – a person can have more than one phone numbers
 - Derived attributes: attributes that can be computed from other attributes
 - *E.g.*, age, given *date_of_birth*
- Domain: the set of permitted values for each attribute

Composite Attributes

- Composite attributes allow us to divided attributes into subparts (other attributes)



Representing Complex Attributes in E-R Diagrams




Agenda

- SQL data definition language (DDL)
- Designing a database
- E-R diagrams
 - **Mapping cardinalities**
 - Primary keys in E-R models
 - Weak entity sets
 - Reduction to relation schemas

Mapping Cardinalities


- Express the **number of entities** to which another entity can be **associated via a relationship set**
 - Most useful in describing binary relationship sets
- For a **binary relationship set** the mapping cardinality must be one of the following types:
 - *One to one*
 - *One to many*
 - *Many to one*
 - *Many to many*

Mapping Cardinalities



(a)

One to one




(b)

One to many


- Note: Some elements in A and B may not be mapped to any elements in the other set

Mapping Cardinalities



(a)

Many to one




(b)

Many to many

- Note: Some elements in A and B may not be mapped to any elements in the other set


Representing Cardinalities in E-R Diagrams

- Express cardinality constraints by drawing either a **directed line (\rightarrow)**, signifying “one,” or an **undirected line ($-$)**, signifying “many,” between the relationship set and the entity set
- One-to-one** relationship between an *instructor* and a *student*:
 - A *student* is associated with **at most one** *instructor* via the relationship *advisor*, and *vice versa*




One-to-Many Relationship

- **One-to-many** relationship between an *instructor* and a *student*
 - An *instructor* is associated with **several (including 0)** *students* via *advisor*
 - A *student* is associated with **at most one** *instructor* via *advisor*




Many-to-One Relationship

- **Many-to-one** relationship between an *instructor* and a *student*
 - An *instructor* is associated with **at most one** *student* via *advisor*
 - A *student* is associated with **several (including 0)** *instructors* via *advisor*




Many-to-Many Relationship

- **Many-to-many** relationship between an *instructor* and a *student*
 - An *instructor* is associated with **several (possibly 0)** *students* via *advisor*
 - A *student* is associated with **several (possibly 0)** *instructors* via *advisor*



Total and Partial Participation

- **Total participation** (*indicated by double line*): every entity in an entity set participates in at least one relationship in the relationship set




participation of *student* in *advisor* relation is total

- *E.g.*, Every *student* must have an associated *instructor*
- **Partial participation**: some entities may not participate in any relationship in the relationship set
 - *E.g.*, Participation of *instructor* in *advisor* is partial

Notation for Expressing More Complex Constraints

- A line may have an **associated minimum and maximum cardinality**, shown in the form $l..h$, where l is the minimum and h the maximum cardinality
 - A minimum value of 1 indicates **total participation**
 - A maximum value of 1 indicates that the entity participates in **at most one** relationship
 - A maximum value of * indicates **no limit**
- Examples
 - Instructor can advise 0 or more students
 - A student must have 1 advisor; cannot have multiple advisors



Agenda

- SQL data definition language (DDL)
- Designing a database
- E-R diagrams
 - Mapping cardinalities
 - **Primary keys in E-R models**
 - Weak entity sets
 - Reduction to relation schemas

Primary Key

- Primary keys provide a way to specify **how entities and relationships are distinguished**
- We consider:
 - Entity sets
 - Relationship sets
 - Weak entity sets

Primary Key for Entity Sets

- By definition, individual entities are distinct
- From database perspective, the *differences among entities must be expressed in terms of their attributes*
 - The attribute values of an entity must be such that *they can uniquely identify the entity*
 - No two entities in an entity set are allowed to have exactly the same value for all attributes
- A *key* for an entity is a set of attributes that *suffice to distinguish entities* from each other

Primary Key for Relationship Sets

- To distinguish among the various relationships of a relationship set, use the **individual primary keys of the entities** in the relationship set
 - Let R be a relationship set involving entity sets E_1, E_2, \dots, E_n
 - The primary key for R consists of the **union of the primary keys** of entity sets E_1, E_2, \dots, E_n
 - If the relationship set R has attributes a_1, a_2, \dots, a_m associated with it, then the primary key of R also includes the attributes a_1, a_2, \dots, a_m
- Example: relationship set “*advisor*”
 - The primary key consists of *instructor.ID* and *student.ID*

Choice of Primary Key for Binary Relationship

- The choice of the primary key for a relationship set **depends on the mapping cardinality** of the relationship set
 - **Many-to-Many** relationships: The preceding **union of the primary keys** is a minimal super key and is chosen as the primary key
 - **One-to-Many** relationships: The **primary key of the “Many” side** is a minimal super key and is used as the primary key
 - **Many-to-one** relationships: The **primary key of the “Many” side** is a minimal super key and is used as the primary key
 - **One-to-one** relationships: The **primary key of either one of the participating entity sets** forms a minimal super key, and either one can be chosen as the primary key

Agenda

- SQL data definition language (DDL)
- Designing a database
- E-R diagrams
 - Mapping cardinalities
 - Primary keys in E-R models
 - **Weak entity sets**
 - Reduction to relation schemas

Weak Entity Sets

- A **weak entity set** is one whose existence is **dependent on another entity**, called its **identifying entity**
- Instead of associating a primary key with a weak entity, **use the identifying entity**, along with extra attributes called **discriminator** to uniquely identify a weak entity
 - A **weak entity** set **does not have a primary key**
 - We still need a *means of distinguishing* among an entity set
 - **Discriminator of a weak entity**: a set of attributes allowing such distinction
 - **Primary key of a weak entity set**
 - = **primary key of a strong entity set** (which its existence depends)
+ its **discriminator**

Weak Entity Sets

- A **weak entity set** is one whose existence is **dependent on another** entity, called its **identifying entity**
- Instead of associating a primary key with a weak entity, **use the identifying entity**, along with extra attributes called **discriminator** to uniquely identify a weak entity
- *E.g.*, Consider a *section* entity, which is uniquely identified by a *course_id*, *semester*, *year*, and *sec_id* → Section entities are related to course entities
 - Treat the relationship *sec_course* as a special relationship that provides extra information
 - In this case, the *course_id*, required to identify *section* entities uniquely



Weak Entity Sets

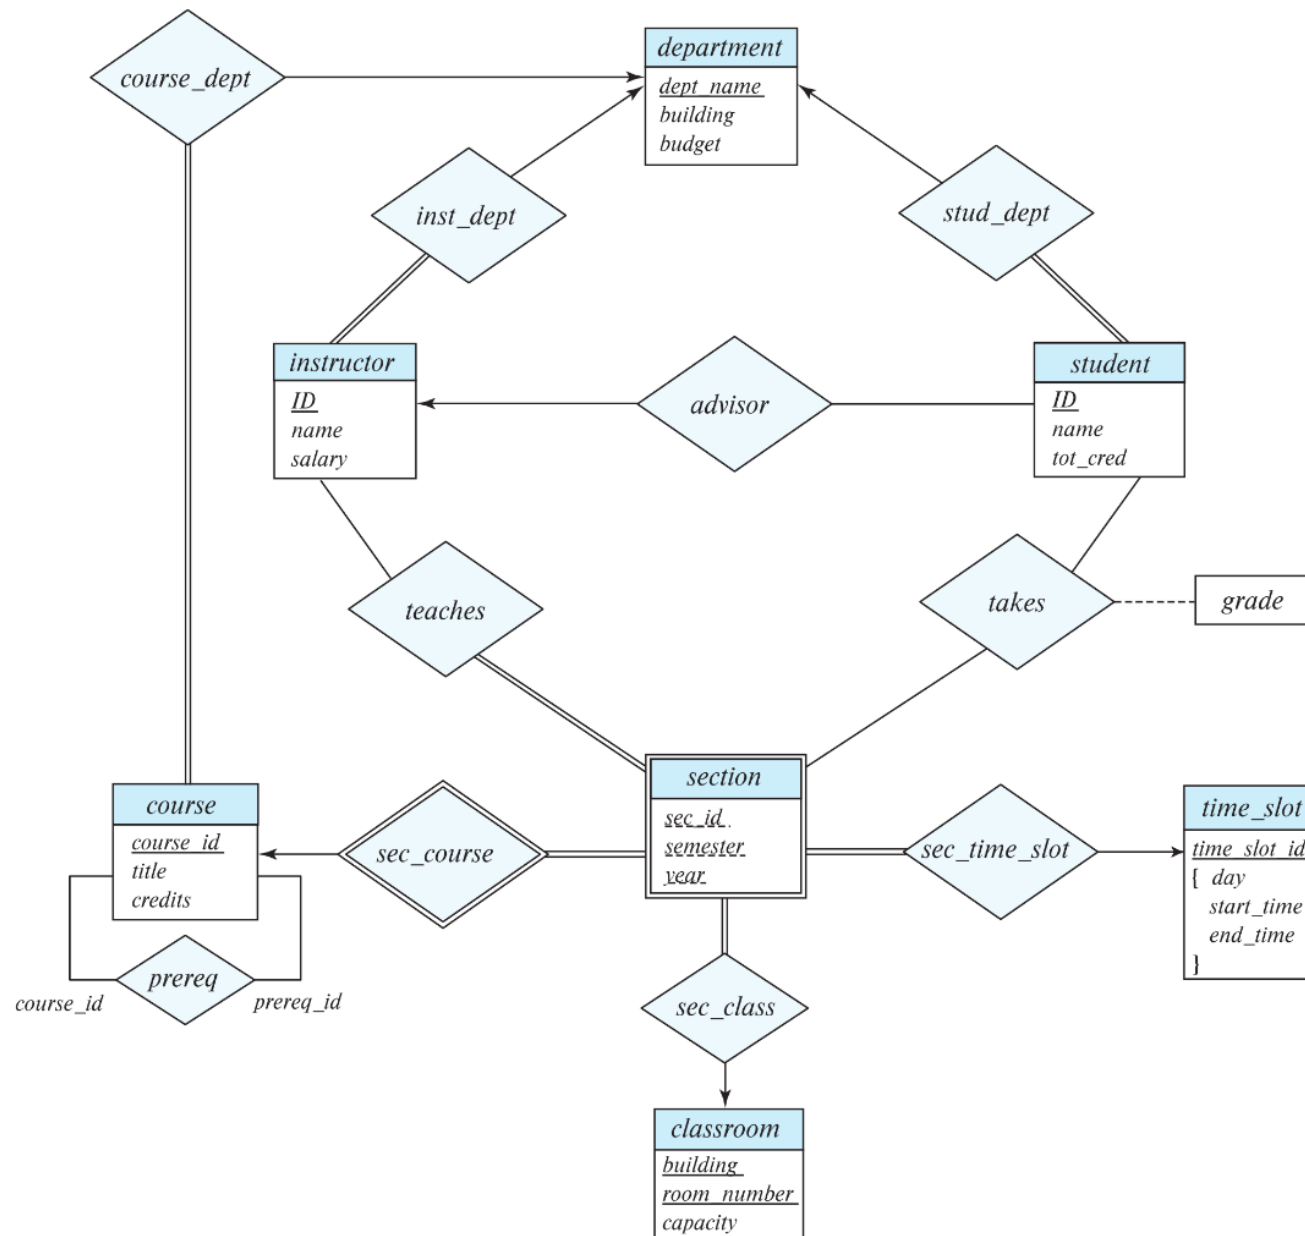
- Identifying entity
 - Every weak entity must be associated with an identifying entity;
 - That is, the weak entity set is said to be **existence dependent** on the identifying entity set
- The **identifying entity set** is said to **own** the weak entity set that it identifies
 - Identifying entity set: an entity set that **has a primary key**
 - Identifying entity set = **strong entity set**
- Identifying relationship
 - **Identifying relationship**: The relationship associating the weak entity set with the identifying entity set

Expressing Weak Entity Sets

- A **weak entity set** is depicted via a double rectangle
- Underline the **discriminator** of a weak entity set with a dashed line
- The **relationship** set connecting the weak entity set to the identifying strong entity set is depicted by a double diamond
 - E.g., Primary key for section – (*course_id*, *sec_id*, *semester*, *year*)



E-R Diagram for a *University* Database



Example: a Record Shop

- Entity sets
 - *customer*
 - *product* (CD, vinyl)
 - *purchase*
- Relationship sets
 - *contains* (between *product* and *purchase*)
 - *buyer* (between *customer* and *purchase*)
- An entity is a specific object; *E.g.*, a newest CD of BTS
- A relationship is specific pair of related objects

Example: a Record Shop

- Attributes
 - *customer*
 - *name, address, phone number, email, etc.*
 - *product*
 - *artist, title, price, description*
 - *purchase*
 - *date, payment_method*
- Artificial primary keys for all entity sets

Example: a Record Shop

- Types of Relationships
 - 1-to-1
 - 1-to-many
 - *buyer* relation (between *customer* and *purchase*)
 - Many-to-many
 - *contains* relation (between *purchase* and *product*)

E-R Diagram

Example: Flight Database

- Entity sets
 - *airport*
 - *code, name, city*
 - *flight*
 - *flight_num, STD, STA, date_offset*
 - *departure*
 - *dept_date, capacity*
 - *customer*
 - *id, name, address, milege_num*
- Relationship sets
 - *to, from (flight, airport)*
 - *flight_dept (flight, departure)*
 - *reserved_on (customer, departure)*

E-R Diagram

Agenda

- E-R diagrams
 - Mapping cardinalities
 - Primary keys in E-R models
 - Weak entity sets
 - **Reduction to relation schemas**

Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as **relation schemas**
 - For each entity set and relationship set, there is a unique schema that is assigned the name of the corresponding entity set or relationship set
 - Each schema has a number of columns (generally corresponding to attributes), which have unique names

Representing Entity Sets

- A **strong entity set** reduces to a **schema with the same attributes**
 - *E.g., student(ID, name, tot_cred)*
- A **weak entity set** becomes a table that includes a **column for the primary key of the identifying strong entity set**
 - *E.g., section (course_id, sec_id, sem, year)*

Representation of Entity Sets with Composite Attributes

<i>instructor</i>
<u>ID</u>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
<i>city</i>
<i>state</i>
<i>zip</i>
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age</i> ()


- Composite attributes are **flattened out** by creating a separate attribute for each component attribute
 - E.g., *first_name* → *name_first_name*
 last_name → *name_last_name*
 - Prefixes can be omitted if there is no ambiguity
- E.g., Ignoring multivalued attributes (*phone_number*), a corresponding instructor schema is:
 - *instructor*(ID,
 first_name, *middle_initial*, *last_name*,
 street_number, *street_name*, *apt_number*,
 city, *state*, *zip_code*,
 date_of_birth)

Representation of Entity Sets with Multivalued Attributes

- A multivalued attribute M of an entity E is represented by a **separate schema EM**
 - Schema EM has attributes corresponding to the **primary key of E** and **an attribute corresponding to multivalued attribute M**
 - *E.g.*, Multivalued attribute *phone_number* of *instructor*:
 - *inst_phone*(ID, phone_number)
- Each value of the multivalued attribute maps to a separate tuple of the relation on schema EM
 - *E.g.*, an *instructor* entity with primary key 22222 and phone numbers 456-7890 and 123-4567
→ maps to two tuples


Representing Relationship Sets

- Any relationship set of strong entity sets can be represented as a schema with attributes for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set
 - E.g., schema for relationship set *advisor*
 - $advisor = (\underline{s_id}, \underline{i_id})$



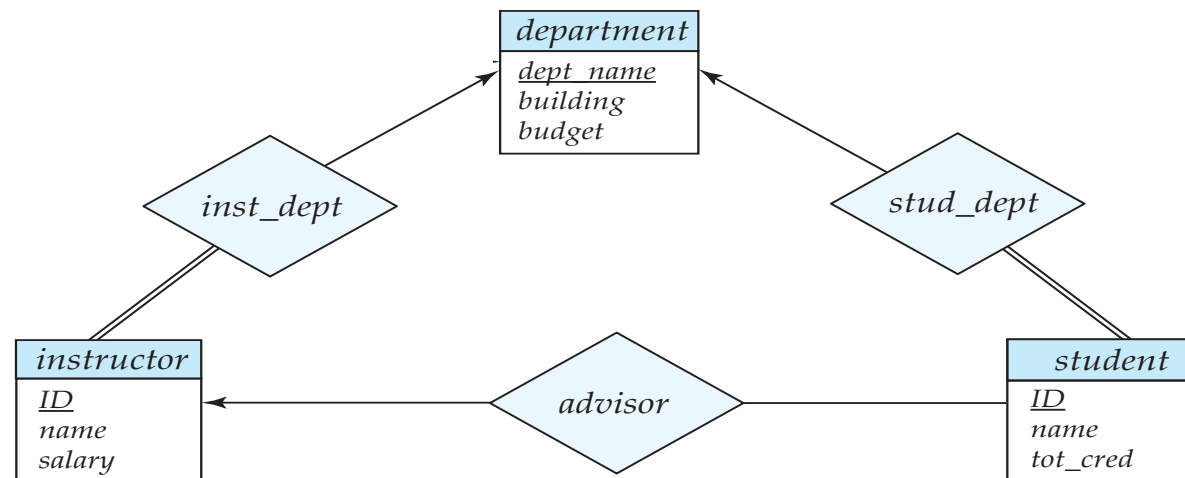
Representing Relationship Sets

- Any relationship set of strong entity sets can be represented as a schema with attributes for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set
 - E.g., schema for relationship set *advisor*
 - $advisor = (\underline{s_id}, \underline{i_id})$



Redundancy of Schemas

- Such “mapping tables” may be redundant
 - Many-to-one and one-to-many relationship sets that are total on the many-side
 - Can be represented by adding an extra attribute to the “many” side, containing the primary key of the “one” side
 - E.g., Instead of creating a schema for relationship set *inst_dept*, add an attribute *dept_name* to the schema arising from entity set *instructor*



Redundancy of Schemas

- Such “mapping tables” may be redundant
 - Many-to-one and one-to-many relationship sets that are total on the many-side
 - Can be represented by adding an extra attribute to the “many” side, containing the primary key of the “one” side
 - E.g., Instead of creating a schema for relationship set *inst_dept*, add an attribute *dept_name* to the schema arising from entity set *instructor*
 - When participation is partial on the “many” side, replacing a schema by an extra attribute in the schema corresponding to the “many” side could result in null values

Redundancy of Schemas

- Such “mapping tables” may be redundant
 - For **one-to-one** relationship sets, either side can be chosen to act as the “many” side
 - An extra attribute can be added to either of the tables corresponding to the two entity sets

EOF

- Coming next:
 - Normalization theory