

ECE30030/ITP30010 Database Systems

SQL DDL

Reading: Chapter 3

Charmgil Hong

charmgil@handong.edu

Spring, 2025

Handong Global University



Announcements

- HW#2 is released
 - Due: Thursday, April 10
- Quiz #1 is on April 3, 2025 (this coming Thursday)
 - Regular class hours (at OH401)
 - Section 1: 2:30pm
 - Section 2: 4pm
 - Closed-book
 - Topics
 - SELECT, FROM, WHERE
 - NULL values
 - Set operations
 - String operations, ordering
 - Aggregate functions, aggregation

Announcements

- Office hours
 - Usually Thursday 5:30pm, Friday 1:30pm
 - Please use Calendly to make an appointment
 - <https://calendly.com/charmgil/sched>
 - Office: NTH201
- Teaching assistants
 - 박동건 systec24@handong.ac.kr
 - 안정현 daniel7806@handong.ac.kr
 - 정광진 21900651@handong.ac.kr
 - 최윤성 22000758@handong.ac.kr

SQL Order of Execution

Order	Clause	Function
1	FROM	Choose and join tables to get base data
2	WHERE	Filters the base data
3	GROUP BY	Aggregates the base data
4	HAVING	Filters the aggregated data
5	SELECT	Returns the final data
6	ORDER BY	Sorts the final data
7	LIMIT	Limits the returned data to a row count

```
SELECT developer, COUNT(*) AS GamesCount, AVG(Price) AS AvgPrice
FROM games
JOIN distributors ON games.dist_id = distributors.dist_id
WHERE release_year > 2019
GROUP BY developer
HAVING COUNT(*) > 2
ORDER BY AvgPrice DESC;
```


SQL Order of Execution

- Explanation
 1. **FROM** games: The data source is defined
 2. **JOIN** distributors **ON** games.dist_id = distributors.dist_id: Tables are joined to include distributor data in the results
 3. **WHERE** release_year > 2019: Rows are filtered based on the release year of the games
 4. **GROUP BY** developer: The results are grouped by the developer
 5. **HAVING COUNT(*) > 2**: The groups are filtered based on the count of games per developer
 6. **SELECT** developer, **COUNT(*)** AS GamesCount, **AVG(Price)** **AS** AvgPrice: Specific columns and calculated columns are selected
 7. **ORDER BY** AvgPrice **DESC**: The final result set is ordered by the average price in descending order

Agenda

- SQL DDL (Data Definition Language)

SQL Commands



Data Definition Language

- The SQL data-definition language (DDL) allows the specification of information about relations, including:
 - The **schema** for each relation
 - The **type** of values associated with each attribute
 - The Integrity **constraints**
 - The set of **indices** to be maintained for each relation
 - Security and authorization information for each relation
 - The physical storage structure of each relation on disk

CREATE DATABASE

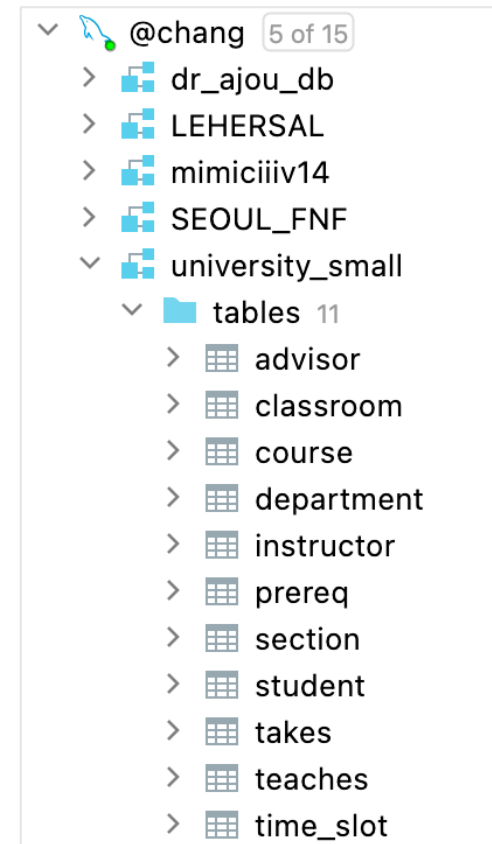
- To initialize a new database
- Basic syntax:
CREATE DATABASE *database_name*
- One can specify the default character encoding method along with this command

- **CREATE DATABASE** test
DEFAULT CHARACTER SET utf8
COLLATE utf8_unicode_ci;

- Collation: a set of rules that defines how to compare and sort character strings

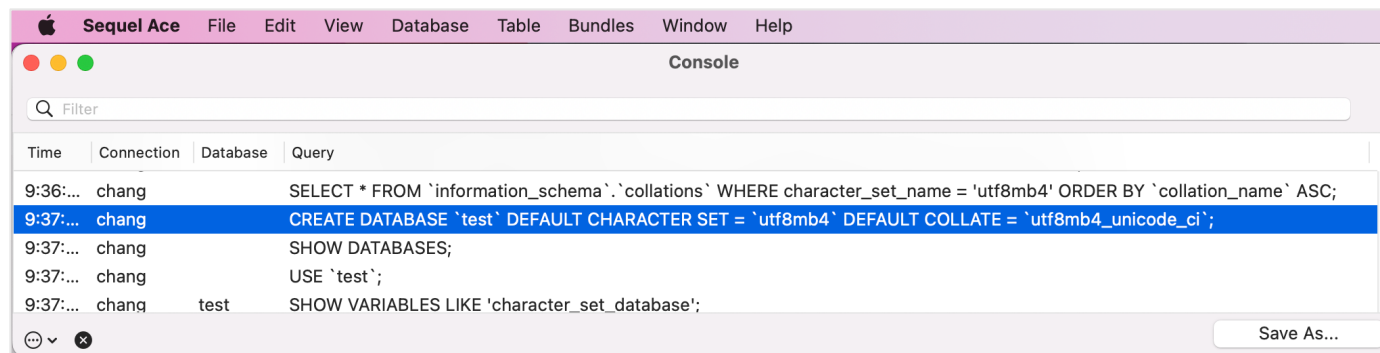
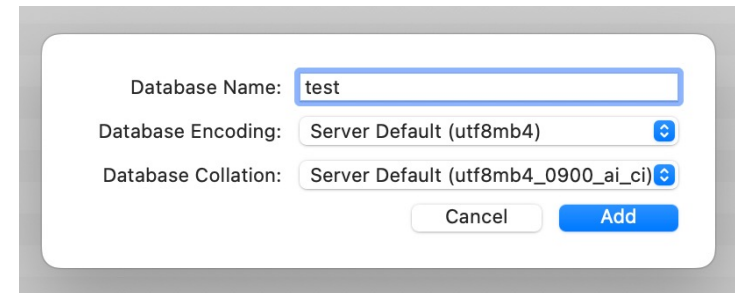
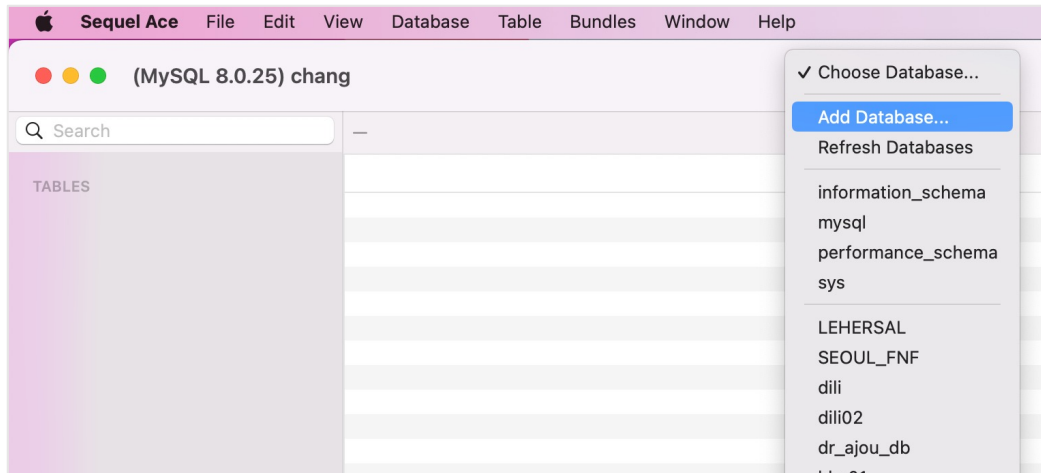
- After creating a database, to use it
USE *database_name*

- Ref: <https://dev.mysql.com/doc/refman/8.4/en/charset.html>



Example: Creating a Database on Sequel Pro

- Creating a new database



CREATE TABLE

- To create a new table

- Basic syntax:

```
CREATE TABLE table_name(  
    Col1_name      data_type[(size)],  
    Col2_name      data_type[(size)]  
)
```

- E.g., Creating a table with four columns

- **CREATE TABLE** *books*(
 ISBN CHAR(20),
 Title CHAR(50),
 AuthorID INTEGER,
 Price FLOAT)

C.f., Table Updates (Updating Tuples)

- INSERT
 - **INSERT INTO** *instructor* **VALUES** ('10211', 'Smith', 'Biology', 66000)
- DELETE
 - **DELETE FROM** *student*
 - Remove all tuples from the *student* relation

Table Updates (Updating Table Schemas)

- DROP TABLE
 - **DROP TABLE r**
 - Remove relation r
- ALTER
 - **ALTER TABLE r ADD A D**
 - A is the name of the **new attribute** to add to relation r ; D is the **domain** of A
 - All **existing tuples in the relation are assigned *null*** as the value for the new attribute
 - **ALTER TABLE r DROP A**
 - A is the name of an **attribute** in r
 - Dropping of attributes not supported by many databases (MySQL does)

Table Updates (Updating Table Schemas)

- Examples
 - **DROP TABLE** time_slot_backup;
 - **ALTER TABLE** time_slot_backup **ADD** remark VARCHAR(20);
 - **ALTER TABLE** time_slot_backup **DROP** remark;

Data Definition Language (revisited)

- The SQL data-definition language (DDL) allows the specification of information about relations, including:
 - The **schema** for each relation
 - The **type** of values associated with each attribute
 - The Integrity **constraints**
 - The set of **indices** to be maintained for each relation
 - Security and authorization information for each relation
 - The physical storage structure of each relation on disk

Data Types in SQL

- Following categories of data types exist in most DBMSs
 - String data
 - Numeric data
 - Temporal data
 - Large objects

Domain Types in SQL

- SQL Data Types
 - **CHAR(*n*)**: Fixed length character string, with user-specified length *n*
 - Maximum length *n* = [0, 255]
 - **VARCHAR(*n*)**: Variable length character strings, with user-specified maximum length *n*
 - Maximum length *n* = [0, 65,535]
 - *If the length is always the same, use a CHAR-type attribute;
if you are storing wildly variable length strings, use a VARCHAR-type attribute*
 - **TEXT**: for strings longer than the range of VARCHAR
 - TINYTEXT 0 – 255 bytes
 - TEXT 0 – 65,535 bytes
 - MEDIUMTEXT 0 – 16,777,215 bytes
 - LONGTEXT 0 – 4,294,967,295 bytes

String Data in SQL

- Difference between CHAR and VARCHAR

Value	CHAR(4)	Storage	VARCHAR(4)	Storage
''	' '	4 bytes	''	1 bytes
'ab'	'ab '	4 bytes	'ab'	3 bytes
'abcd'	'abcd'	4 bytes	'abcd'	5 bytes
'abcdefg'	'abcd'	4 bytes	'abcd'	5 bytes

Domain Types in SQL

- `"\"%ab%\""`

Domain Types in SQL

- SQL Data Types
 - **INT, INTEGER**: **Integer** (a finite subset of the integers that is machine-dependent)
 - **SMALLINT**: **Small integer** (a machine-dependent subset of the integer domain type)
 - **BIGINT**: **Big integer** (a machine-dependent subset of the integer domain type)
- **TINYINT** and **MEDIUMINT** are also available

Domain Types in SQL

- Different R-DBMSs support different combinations of those integer types

	Bytes	MySQL	MS SQL	PostgreSQL	DB2
TINYINT	1	✓	✓		
SMALLINT	2	✓	✓	✓	✓
MEDIUMINT	3	✓			
INT/INTEGER	4	✓	✓	✓	✓
BIGINT	8	✓	✓	✓	✓

- *C.f.*, Oracle only has a NUMBER datatype

Domain Types in SQL

- SQL Data Types
 - **NUMERIC(p,d)**: **Fixed point number** (exact value) with user-specified precision of p digits, with d digits to the right of decimal point
 - *E.g.*, **NUMERIC**(3,1) allows 44.5 to be stores exactly, but not 444.5 or 0.32)
 - In MySQL, **DECIMAL** is NUMERIC
 - **FLOAT**: Floating point number (approximate) with **single-precision**
 - **REAL, DOUBLE**: Floating point number (approximate) with **double-precision**

Domain Types in SQL

- DECIMAL vs INT/FLOAT/DOUBLE
 - FLOAT and DOUBLE are faster than DECIMAL
 - DECIMAL values are exact
 - Example

floats: FLOAT	decimals: DECIMAL(3,2)
1.1	1.10
1.1	1.10
1.1	1.10

- SELECT SUM(...) → **DECIMAL values are precise**

SUM(floats)	SUM(decimals)
3.3000000715255737	3.30

Domain Types in SQL

- SQL Data Types
 - **DATE**: 'YYYY-MM-DD'
 - Range: 1000-01-01 to 9999-12-31
 - *E.g.*, '2020-03-01' for March 1, 2020
 - **TIME**: 'HH:MM:SS'
 - Range: -838:59:59 to 838:59:59
 - *E.g.*, '14:30:03.5' for 3.5 seconds after 2:30pm
 - **DATETIME**: 'YYYY-MM-DD HH:MM:SS'
 - Range: 1000-01-01 00:00:00 to 9999-12-31 23:59:59
 - **YEAR**: 'YYYY'
 - Range: 1901 to 2155, or 0000 (illegal year values are converted to 0000)

Temporal Data in SQL

- SQL Data Types
 - **TIMESTAMP(*n*)**: **Unix time** (time since Jan 1, 1970)
 - A way to track time as a running total of seconds
 - Range: 1970-01-01 00:00:01 UTC to 2038-01-19 03:14:07 UTC
 - Typically used for logging (keeping records of all the system events)
 - URL: <https://time.is/Unix>

**UNIX
TIME
SINCE 1970**

Temporal Data in SQL

- SQL Data Types
 - **TIMESTAMP(*n*)**: **Unix time** (time since Jan 1, 1970)
 - A way to track time as a running total of seconds
 - Range: 1970-01-01 00:00:01 UTC to 2038-01-19 03:14:07 UTC
 - Typically used for logging (keeping records of all the system events)
 - URL: <https://time.is/Unix>

Binary : 01111111 11111111 11111111 11110000

Decimal : 2147483632

Date : 2038-01-19 03:13:52 (UTC)

Date : 2038-01-19 03:13:52 (UTC)



Domain Types in SQL

- SQL Data Types
 - **TIMESTAMP(*n*)**: **Unix time** (time since Jan 1, 1970)
 - Range: 1970-01-01 00:00:01 UTC to 2038-01-19 03:14:07 UTC
 - Typically used for logging (keeping records of all the system events)
 - Depending on size *n*, the display pattern changes

	Format
TIMESTAMP(14)	YYYYMMDDHHMMSS
TIMESTAMP(12)	YYMMDDHHMMSS
TIMESTAMP(10)	YYMMDDHHMM
TIMESTAMP(8)	YYYYMMDD
TIMESTAMP(6)	YYMMDD
TIMESTAMP(4)	YYMM
TIMESTAMP(2)	YY

Domain Types in SQL

- SQL Data Types
 - **BINARY(*n*)**: binary byte data type, with user-specified length *n*
 - Contains a byte strings (rather than a character string)
 - Maximum length $n = [0, 255]$
 - **VARBINARY(*n*)**: binary byte data type, with user-specified maximum length *n*
 - Maximum length $n = [0, 65,535]$
 - **BLOB**: Binary Large OBject data type
 - TINYBLOB 0 – 255 bytes
 - BLOB 0 – 65,535 bytes (65 KB)
 - MEDIUMBLOB 0 – 16,777,215 bytes (16 MB)
 - LONGBLOB 0 – 4,294,967,295 bytes (4 GB)

CREATE TABLE Construct

- A new relation is defined using the **CREATE TABLE** command:

CREATE TABLE *r*

(*A*₁ *D*₁, *A*₂ *D*₂, ..., *A*_{*n*} *D*_{*n*},
(*integrity-constraint*₁),
...,
(*integrity-constraint*_{*k*}))

- *r* is the name of the relation
 - Each *A*_{*i*} is an attribute name in the schema of relation *r*
 - Each *D*_{*i*} is the data type of values in the domain of attribute *A*_{*i*}
-
- Example: **CREATE TABLE** instructor(

ID	CHAR(5),
name	VARCHAR(20),
dept_name	VARCHAR(20),
salary	NUMERIC(8,2))

Integrity Constraints in CREATE TABLE

- SQL prevents any update to the database that violates an **integrity constraint**
 - Integrity constraints allow us to specify what data makes sense for us
- Types of integrity constraints
 - Primary key: **PRIMARY KEY** (A_1, \dots, A_n)
 - Foreign key: **FOREIGN KEY** (A_m, \dots, A_n) **REFERENCES** r
 - Unique key: **UNIQUE**
 - Not null: **NOT NULL**

- Example:

```
CREATE TABLE instructor(  
    ID                CHAR(5),  
    name              VARCHAR(20) NOT NULL,  
    dept_name         VARCHAR(20)  
    salary           NUMERIC(8, 2),  
    PRIMARY KEY (ID),  
    FOREIGN KEY (dept_name) REFERENCES department);
```

Declaring Keys

- An attribute or list of attributes may be declared as PRIMARY KEY or UNIQUE
 - Meaning: no two tuples of the relation may agree in all the attribute(s) on the list
 - That is, the attribute(s) do(es) **not allow duplicates** in values
 - PRIMARY KEY/UNIQUE can be used as an **identifier for each row**
 - Comparison: PRIMARY KEY vs UNIQUE

PRIMARY KEY	UNIQUE
Used to serve as a unique identifier for each row in a relation	Uniquely determines a row which is not primary key
Cannot accept NULL	Can accept NULL values (some DBMSs accept only one NULL value)
A relation can have only one primary key	A relation can have more than one unique attributes
Clustered index	Non-clustered index

Integrity Constraints

- **NOT NULL** – disallowing null values
 - Null values indicate that the data is not known
 - These can cause problems in querying database
 - The Primary Key columns automatically prevent null being entered
 - *C.f.*, **NULL** – can be used to **explicitly allow** null values

```
CREATE TABLE studio (  
    ID                NUMERIC(5,0) PRIMARY KEY,  
    name             VARCHAR(20) NOT NULL,  
    city             VARCHAR(20) NULL,  
    state            CHAR(2) NOT NULL  
);
```


Integrity Constraints

- **DEFAULT** – A default value can be inserted in any column with this keyword

- *E.g.*, **CREATE TABLE** *movies*(
 movie_title **VARCHAR(40) NOT NULL,**
 release_date **DATE DEFAULT sysdate NULL,**
 genre **VARCHAR(20) DEFAULT 'Comedy'**
 CHECK genre IN ('Comedy', 'Action', 'Drama')
)

- In MySQL,
 - **CREATE TABLE** *movies*(
 movie_title **VARCHAR(40) NOT NULL,**
 release_date **DATE DEFAULT CURRENT_TIMESTAMP NULL,**
 genre **VARCHAR(20) DEFAULT 'Comedy'**
 CHECK genre IN ('Comedy', 'Action', 'Drama')
)

Integrity Constraints

- **CHECK** – Allows the inserted value to be checked
 - *E.g.*, **CREATE TABLE** *movies*(
 movie_title **VARCHAR(40) PRIMARY KEY,**
 release_date **DATE,**
 budget **INTEGER CHECK (budget > 50000)**
)
 - Table-level constraints can be defined; *E.g.*,
 - **CREATE TABLE** *movies*(
 movie_title **VARCHAR(40) PRIMARY KEY,**
 release_date **DATE,**
 budget **INTEGER CHECK (budget > 50000),**
 CONSTRAINT *release_date_const*
 CHECK (release_date BETWEEN '01-Jan-2000' AND '31-Dec-2009')
)

Declaring Keys

- **CREATE TABLE** *student* (
 ID **VARCHAR(5),**
 name **VARCHAR(20) NOT NULL,**
 dept_name **VARCHAR(20),**
 tot_cred **NUMERIC(3,0),**
 PRIMARY KEY (ID),
 FOREIGN KEY (dept_name) REFERENCES department);

- **CREATE TABLE** *student* (
 ID **VARCHAR(5) PRIMARY KEY,**
 name **VARCHAR(20) NOT NULL,**
 dept_name **VARCHAR(20),**
 tot_cred **NUMERIC(3,0),**
 FOREIGN KEY (dept_name) REFERENCES department);

More Examples

- **CREATE TABLE** *takes* (
 ID **VARCHAR**(5),
 course_id **VARCHAR**(8),
 sec_id **VARCHAR**(8),
 semester **VARCHAR**(6),
 year **NUMERIC**(4,0),
 grade **VARCHAR**(2),
 PRIMARY KEY (*ID*, *course_id*, *sec_id*, *semester*, *year*),
 FOREIGN KEY (*ID*) **REFERENCES** *student*,
 FOREIGN KEY (*course_id*, *sec_id*, *semester*, *year*)
 REFERENCES *section*);

More Examples

- **CREATE TABLE** *course* (
 course_id **VARCHAR**(8),
 title **VARCHAR**(50),
 dept_name **VARCHAR**(20) **DEFAULT** 'Comp. Sci',
 credits **NUMERIC**(2,0),
 PRIMARY KEY (*course_id*),
 FOREIGN KEY (*dept_name*) **REFERENCES** *department*);

More Examples

- **CREATE TABLE** *neighbors*(
 name **CHAR(30) PRIMARY KEY**,
 addr **CHAR(50) DEFAULT '123 Sesame St.'**,
 phone **CHAR(16));**
- Inserting Elmo is a neighbor:
 - **INSERT INTO** *neighbors* (*name*)
 VALUES ('Elmo');

name	addr	phone
'Elmo'	'123 Sesame St.'	NULL

More Examples

- **CREATE TABLE** *neighbors*(
 name **CHAR(30) PRIMARY KEY**,
 addr **CHAR(50) DEFAULT '123 Sesame St.'**,
 phone **CHAR(16) NOT NULL**);
- Inserting Elmo is a neighbor:
 - **INSERT INTO** *neighbors* (*name*)
 VALUES ('Elmo');
 - ➔ If *phone* were NOT NULL, this insertion would have been **rejected**